

projects/project_1/gan/gan.py

```
import numpy as np
```

```
import torch
```

```
import torch.nn as nn
```

```
from torch.utils.data import sampler
```

```
NOISE_DIM = 96
```

```
dtype = torch.cuda.FloatTensor if torch.cuda.is_available() else torch.FloatTensor
```

```
def sample_noise(batch_size, dim, seed=None):
```

```
    """
```

```
    Generate a PyTorch Tensor of uniform random noise.
```

```
    Input:
```

- batch_size: Integer giving the batch size of noise to generate.
- dim: Integer giving the dimension of noise to generate.

```
    Output:
```

- A PyTorch Tensor of shape (batch_size, dim) containing uniform random noise in the range (-1, 1).

```
    """
```

```
    if seed is not None:
```

```
        torch.manual_seed(seed)
```

```
    # *****START OF YOUR CODE (DO NOT DELETE/MODIFY THIS LINE)*****
```

```
    return torch.FloatTensor(batch_size, dim).uniform_(-1, 1)
```

```
    # *****END OF YOUR CODE (DO NOT DELETE/MODIFY THIS LINE)*****
```

```
def discriminator(seed=None):
```

```
    """
```

```
    Build and return a PyTorch model implementing the architecture above.
```

```
    """
```

```
    if seed is not None:
```

```
        torch.manual_seed(seed)
```

```
    model = None
```

```
#####
```

```
# TODO: Implement architecture #
```

```
#
```

```
#
```

```
# HINT: nn.Sequential might be helpful. You'll start by calling Flatten(). #
```

```
#####
```

```
# *****START OF YOUR CODE (DO NOT DELETE/MODIFY THIS LINE)*****
```

```

model = nn.Sequential(
    nn.Flatten(),
    nn.Linear(784, 256),
    nn.LeakyReLU(negative_slope=0.01),
    nn.Linear(256, 256),
    nn.LeakyReLU(negative_slope=0.01),
    nn.Linear(256, 1),
)
pass

# *****END OF YOUR CODE (DO NOT DELETE/MODIFY THIS LINE)*****
#####
#                                     END OF YOUR CODE                                     #
#####
return model

def generator(noise_dim=NOISE_DIM, seed=None):
    """
    Build and return a PyTorch model implementing the architecture above.
    """

    if seed is not None:
        torch.manual_seed(seed)

    model = None

    #####
    # TODO: Implement architecture                                                    #
    #                                                                              #
    # HINT: nn.Sequential might be helpful.                                         #
    #####
    # *****START OF YOUR CODE (DO NOT DELETE/MODIFY THIS LINE)*****

    model = nn.Sequential(
        nn.Linear(noise_dim, 1024),
        nn.ReLU(),
        nn.Linear(1024, 1024),
        nn.ReLU(),
        nn.Linear(1024, 784),
        nn.Tanh(),
    )

    # *****END OF YOUR CODE (DO NOT DELETE/MODIFY THIS LINE)*****
    #####
    #                                     END OF YOUR CODE                                     #
    #####
    return model

```

```
def bce_loss(input, target):
    """
    Numerically stable version of the binary cross-entropy loss function in PyTorch.
    Inputs:
    - input: PyTorch Tensor of shape (N, ) giving scores.
    - target: PyTorch Tensor of shape (N,) containing 0 and 1 giving targets.
      dtype is float! (a global dtype is defined above).
    Returns:
    - A PyTorch Tensor containing the mean BCE loss over the minibatch of input data.
    """
    bce = nn.BCEWithLogitsLoss()
    return bce(input, target)
```

```
def discriminator_loss(logits_real, logits_fake):
    """
    Computes the discriminator loss described above.
    Inputs:
    - logits_real: PyTorch Tensor of shape (N,) giving scores for the real data.
    - logits_fake: PyTorch Tensor of shape (N,) giving scores for the fake data.
    Returns:
    - loss: PyTorch Tensor containing (scalar) the loss for the discriminator.
    """
    loss = None
    # *****START OF YOUR CODE (DO NOT DELETE/MODIFY THIS LINE)*****

    true_labels = torch.ones_like(logits_real) # .type(dtype)
    real_loss = bce_loss(logits_real, true_labels)

    fake_labels = torch.zeros_like(logits_fake) # .type(dtype)
    fake_loss = bce_loss(logits_fake, fake_labels)

    loss = real_loss + fake_loss

    # *****END OF YOUR CODE (DO NOT DELETE/MODIFY THIS LINE)*****
    return loss
```

```
def generator_loss(logits_fake):
    """
    Computes the generator loss described above.
    Inputs:
    - logits_fake: PyTorch Tensor of shape (N,) giving scores for the fake data.
    Returns:
    - loss: PyTorch Tensor containing the (scalar) loss for the generator.
    """
    loss = None
    # *****START OF YOUR CODE (DO NOT DELETE/MODIFY THIS LINE)*****
```

```

    targets = torch.ones_like(logits_fake).type(dtype)
    loss = bce_loss(logits_fake, targets)

# *****END OF YOUR CODE (DO NOT DELETE/MODIFY THIS LINE)*****
return loss

def get_optimizer(model):
    """
    Construct and return an Adam optimizer for the model with learning rate 1e-3,
    beta1=0.5, and beta2=0.999.
    Input:
    - model: A PyTorch model that we want to optimize.
    Returns:
    - An Adam optimizer for the model with the desired hyperparameters.
    """
    optimizer = None
    # *****START OF YOUR CODE (DO NOT DELETE/MODIFY THIS LINE)*****

    optimizer = torch.optim.Adam(params=model.parameters(), lr=1e-3, betas=(0.5, 0.999))

    # *****END OF YOUR CODE (DO NOT DELETE/MODIFY THIS LINE)*****
    return optimizer

def ls_discriminator_loss(scores_real, scores_fake):
    """
    Compute the Least-Squares GAN loss for the discriminator.
    Inputs:
    - scores_real: PyTorch Tensor of shape (N,) giving scores for the real data.
    - scores_fake: PyTorch Tensor of shape (N,) giving scores for the fake data.
    Outputs:
    - loss: A PyTorch Tensor containing the loss.
    """
    loss = None
    # *****START OF YOUR CODE (DO NOT DELETE/MODIFY THIS LINE)*****

    real_loss = 0.5 * torch.mean(torch.square(scores_real - 1))
    fake_loss = 0.5 * torch.mean(torch.square(scores_fake))
    loss = real_loss + fake_loss

    # *****END OF YOUR CODE (DO NOT DELETE/MODIFY THIS LINE)*****
    return loss

def ls_generator_loss(scores_fake):
    """
    Computes the Least-Squares GAN loss for the generator.
    Inputs:
    - scores_fake: PyTorch Tensor of shape (N,) giving scores for the fake data.

```

Outputs:

- loss: A PyTorch Tensor containing the loss.

"""

loss = None

*****START OF YOUR CODE (DO NOT DELETE/MODIFY THIS LINE)*****

loss = 0.5 * torch.mean(torch.square(scores_fake - 1))

*****END OF YOUR CODE (DO NOT DELETE/MODIFY THIS LINE)*****

return loss

def build_dc_classifier():

"""

Build and return a PyTorch model for the DCGAN discriminator implementing the architecture above.

"""

#####

TODO: Implement architecture

#

HINT: nn.Sequential might be helpful.

#####

*****START OF YOUR CODE (DO NOT DELETE/MODIFY THIS LINE)*****

```
model = nn.Sequential(
    nn.Conv2d(in_channels=1, out_channels=32, kernel_size=5, stride=1),
    nn.LeakyReLU(negative_slope=0.01),
    nn.MaxPool2d(kernel_size=2, stride=2),
    nn.Conv2d(in_channels=32, out_channels=64, kernel_size=5, stride=1),
    nn.LeakyReLU(negative_slope=0.01),
    nn.MaxPool2d(kernel_size=2, stride=2),
    nn.Flatten(),
    nn.Linear(1024, 1024),
    nn.LeakyReLU(negative_slope=0.01),
    nn.Linear(1024, 1),
)
```

return model

*****END OF YOUR CODE (DO NOT DELETE/MODIFY THIS LINE)*****

#####

END OF YOUR CODE

#####

def build_dc_generator(noise_dim=NOISE_DIM):

"""

Build and return a PyTorch model implementing the DCGAN generator using the architecture described above.

||||

```
#####
# TODO: Implement architecture                                     #
#                                                                 #
# HINT: nn.Sequential might be helpful.                         #
#####
# *****START OF YOUR CODE (DO NOT DELETE/MODIFY THIS LINE)*****
```

```
model = nn.Sequential(
    nn.Linear(noise_dim, 1024),
    nn.ReLU(),
    nn.BatchNorm1d(num_features=1024),
    nn.Linear(1024, 6272),
    nn.ReLU(),
    nn.BatchNorm1d(num_features=6272),
    nn.Unflatten(1, (128, 7, 7)),
    nn.ConvTranspose2d(
        in_channels=128,
        out_channels=64,
        kernel_size=4,
        stride=2,
        padding=1,
    ),
    nn.ReLU(),
    nn.BatchNorm2d(num_features=64),
    nn.ConvTranspose2d(
        in_channels=64,
        out_channels=1,
        kernel_size=4,
        stride=2,
        padding=1,
    ),
    nn.Tanh(),
    nn.Flatten(),
)
```

```
return model
```

```
# *****END OF YOUR CODE (DO NOT DELETE/MODIFY THIS LINE)*****
#####
#                                                                 #
#                               END OF YOUR CODE                  #
#####
```

```
def run_a_gan(
    D,
    G,
    D_solver,
    G_solver,
```

```

discriminator_loss,
generator_loss,
loader_train,
show_every=250,
batch_size=128,
noise_size=96,
num_epochs=10,
):
    """
    Train a GAN!
    Inputs:
    - D, G: PyTorch models for the discriminator and generator
    - D_solver, G_solver: torch.optim Optimizers to use for training the
      discriminator and generator.
    - discriminator_loss, generator_loss: Functions to use for computing the generator
and
      discriminator loss, respectively.
    - show_every: Show samples after every show_every iterations.
    - batch_size: Batch size to use for training.
    - noise_size: Dimension of the noise to use as input to the generator.
    - num_epochs: Number of epochs over the training dataset to use for training.
    """
    images = []
    iter_count = 0
    for epoch in range(num_epochs):
        for x, _ in loader_train:
            if len(x) != batch_size:
                continue
            D_solver.zero_grad()
            real_data = x.type(dtype)
            logits_real = D(2 * (real_data - 0.5)).type(dtype)

            g_fake_seed = sample_noise(batch_size, noise_size).type(dtype)
            fake_images = G(g_fake_seed).detach()
            logits_fake = D(fake_images.view(batch_size, 1, 28, 28))

            d_total_error = discriminator_loss(logits_real, logits_fake)
            d_total_error.backward()
            D_solver.step()

            G_solver.zero_grad()
            g_fake_seed = sample_noise(batch_size, noise_size).type(dtype)
            fake_images = G(g_fake_seed)

            gen_logits_fake = D(fake_images.view(batch_size, 1, 28, 28))
            g_error = generator_loss(gen_logits_fake)
            g_error.backward()
            G_solver.step()

            if iter_count % show_every == 0:

```

```

        print(
            "Iter: {}, D: {:.4}, G:{:.4}".format(
                iter_count, d_total_error.item(), g_error.item()
            )
        )
        imgs_numpy = fake_images.data.cpu().numpy()
        images.append(imgs_numpy[0:16])

    iter_count += 1

return images

```

```

class ChunkSampler(sampler.Sampler):
    """Samples elements sequentially from some offset.
    Arguments:
        num_samples: # of desired datapoints
        start: offset where we should start selecting from
    """

    def __init__(self, num_samples, start=0):
        self.num_samples = num_samples
        self.start = start

    def __iter__(self):
        return iter(range(self.start, self.start + self.num_samples))

    def __len__(self):
        return self.num_samples


class Flatten(nn.Module):
    def forward(self, x):
        N, C, H, W = x.size() # read in N, C, H, W
        return x.view(
            N, -1
        ) # "flatten" the C * H * W values into a single vector per image


class Unflatten(nn.Module):
    """
    An Unflatten module receives an input of shape (N, C*H*W) and reshapes it
    to produce an output of shape (N, C, H, W).
    """

    def __init__(self, N=-1, C=128, H=7, W=7):
        super(Unflatten, self).__init__()
        self.N = N
        self.C = C
        self.H = H

```



```
self.W = W

def forward(self, x):
    return x.view(self.N, self.C, self.H, self.W)

def initialize_weights(m):
    if isinstance(m, nn.Linear) or isinstance(m, nn.ConvTranspose2d):
        nn.init.xavier_uniform_(m.weight.data)

def preprocess_img(x):
    return 2 * x - 1.0

def deprocess_img(x):
    return (x + 1.0) / 2.0

def rel_error(x, y):
    return np.max(np.abs(x - y) / (np.maximum(1e-8, np.abs(x) + np.abs(y))))

def count_params(model):
    """Count the number of parameters in the model."""
    param_count = np.sum([np.prod(p.size()) for p in model.parameters()])
    return param_count
```