

projects/project_1/vae/vae.py

```

from __future__ import print_function
from torch import nn
import torch

```

```

device = torch.device("cpu")

```

```

def hello_vae():
    print("Hello from vae.py!")

```

```

class VAE(nn.Module):
    def __init__(self, input_size, latent_size=15):
        super(VAE, self).__init__()
        self.input_size = input_size # H*W
        self.latent_size = latent_size # Z
        self.hidden_dim = 400 # H_d
        self.encoder = None
        self.mu_layer = None
        self.logvar_layer = None
        self.decoder = None

```

```

#####
#      # TODO: Implement the fully-connected encoder architecture described in the notebook.
#
#      # Specifically, self.encoder should be a network that inputs a batch of input images of
#
#      # shape (N, 1, H, W) into a batch of hidden features of shape (N, H_d). Set up
#
#      # self.mu_layer and self.logvar_layer to be a pair of linear layers that map the hidden
#
#      # features into estimates of the mean and log-variance of the posterior over the latent
#
#      # vectors; the mean and log-variance estimates will both be tensors of shape (N, Z).
#
#####

```

```

#####
# Replace "pass" statement with your code
self.encoder = nn.Sequential(
    nn.Flatten(),
    nn.Linear(self.input_size, self.hidden_dim),
    nn.ReLU(),
    nn.Linear(self.hidden_dim, self.hidden_dim),
    nn.ReLU(),
    nn.Linear(self.hidden_dim, self.hidden_dim),
    nn.ReLU(),
)

self.mu_layer = nn.Linear(self.hidden_dim, self.latent_size)
self.logvar_layer = nn.Linear(self.hidden_dim, self.latent_size)

```

```
#####
# TODO: Implement the fully-connected decoder architecture described in the notebook.
#
# Specifically, self.decoder should be a network that inputs a batch of latent vectors
of #
# shape (N, Z) and outputs a tensor of estimated images of shape (N, 1, H, W).
#
#####
# Replace "pass" statement with your code
self.decoder = nn.Sequential(
    nn.Linear(self.latent_size, self.hidden_dim),
    nn.ReLU(),
    nn.Linear(self.hidden_dim, self.hidden_dim),
    nn.ReLU(),
    nn.Linear(self.hidden_dim, self.hidden_dim),
    nn.ReLU(),
    nn.Linear(self.hidden_dim, self.input_size),
    nn.Sigmoid(),
    nn.Unflatten(dim=1, unflattened_size=(1, 28, 28)),
)

#####
#                                     END OF YOUR CODE
#
#####

def forward(self, x):
    """
    Performs forward pass through FC-VAE model by passing image through
    encoder, reparametrize trick, and decoder models
    Inputs:
    - x: Batch of input images of shape (N, 1, H, W)
    Returns:
    - x_hat: Reconstructed input data of shape (N,1,H,W)
    - mu: Matrix representing estimated posterior mu (N, Z), with Z latent space dimension
    - logvar: Matrix representing estimated variance in log-space (N, Z), with Z latent
space dimension
    """
    x_hat = None
    mu = None
    logvar = None

#####
# TODO: Implement the forward pass by following these steps
#
# (1) Pass the input batch through the encoder model to get posterior mu and logvariance
#
# (2) Reparametrize to compute the latent vector z
#
# (3) Pass z through the decoder to reconstruct x
#
```

```
#####
# Replace "pass" statement with your code
output = self.encoder(x)
mu = self.mu_layer(output)
logvar = self.logvar_layer(output)

reparametrized_output = reparametrize(mu=mu, logvar=logvar)
x_hat = self.decoder(reparametrized_output)

#####
#                                     END OF YOUR CODE
#

#####
return x_hat, mu, logvar

class CVAE(nn.Module):
    def __init__(self, input_size, num_classes=10, latent_size=15):
        super(CVAE, self).__init__()
        self.input_size = input_size # H*W
        self.latent_size = latent_size # Z
        self.num_classes = num_classes # K
        self.hidden_dim = None # H_d
        self.encoder = None
        self.mu_layer = None
        self.logvar_layer = None
        self.decoder = None

#####
# TODO: Define a FC encoder as described in the notebook that transforms the image--
after #
# flattening and now adding our one-hot class vector (N, H*W + K)--into a
hidden_dimension # #
# (N, H_d) feature space, and a final two layers that project that feature space
#
# to posterior mu and posterior log-variance estimates of the latent space (N, Z)
#

#####
# Replace "pass" statement with your code
self.hidden_dim = 400

self.encoder = nn.Sequential(
    nn.Linear(self.input_size + self.num_classes, self.hidden_dim),
    nn.ReLU(),
    nn.Linear(self.hidden_dim, self.hidden_dim),
    nn.ReLU(),
    nn.Linear(self.hidden_dim, self.hidden_dim),
    nn.ReLU(),
)
self.mu_layer = nn.Linear(self.hidden_dim, self.latent_size)
```

```

self.logvar_layer = nn.Linear(self.hidden_dim, self.latent_size)

#####
# TODO: Define a fully-connected decoder as described in the notebook that transforms
the #
# latent space (N, Z + K) to the estimated images of shape (N, 1, H, W).
#
#####
# Replace "pass" statement with your code

self.decoder = nn.Sequential(
    nn.Linear(self.latent_size + self.num_classes, self.hidden_dim),
    nn.ReLU(),
    nn.Linear(self.hidden_dim, self.hidden_dim),
    nn.ReLU(),
    nn.Linear(self.hidden_dim, self.hidden_dim),
    nn.ReLU(),
    nn.Linear(self.hidden_dim, self.input_size),
    nn.Sigmoid(),
    nn.Unflatten(dim=1, unflattened_size=(1, 28, 28)),
)

#####
#                                     END OF YOUR CODE
#
#####

def forward(self, x, labels):
    """
    Performs forward pass through FC-CVAE model by passing image through
    encoder, reparametrize trick, and decoder models
    Inputs:
    - x: Input data for this timestep of shape (N, 1, H, W)
    - labels: One hot vector representing the input class (0-9) (N, K)
    Returns:
    - x_hat: Reconstructed input data of shape (N, 1, H, W)
    - mu: Matrix representing estimated posterior mu (N, Z), with Z latent space dimension
    - logvar: Matrix representing estimated variance in log-space (N, Z), with Z latent
space dimension
    """
    x_hat = None
    mu = None
    logvar = None

#####
# TODO: Implement the forward pass by following these steps
#
# (1) Pass the concatenation of input batch and one hot vectors through the encoder
model #
# to get posterior mu and logvariance
#

```

```

# (2) Reparametrize to compute the latent vector z
#
# (3) Pass concatenation of z and one hot vectors through the decoder to reconstruct x
#
#####
# Replace "pass" statement with your code
flattened = torch.flatten(x, start_dim=1)
inputs = torch.cat((flattened, labels), dim=1)
output = self.encoder(inputs)
mu = self.mu_layer(output)
logvar = self.logvar_layer(output)

z = reparametrize(mu=mu, logvar=logvar)
z = torch.cat((z, labels), dim=1)
x_hat = self.decoder(z)

#####
#                                     END OF YOUR CODE
#

#####
    return x_hat, mu, logvar

def reparametrize(mu, logvar):
    """
    Differentiably sample random Gaussian data with specified mean and variance using the
    reparameterization trick.

    Suppose we want to sample a random number z from a Gaussian distribution with mean mu and
    standard deviation sigma, such that we can backpropagate from the z back to mu and sigma.
    We can achieve this by first sampling a random value epsilon from a standard Gaussian
    distribution with zero mean and unit variance, then setting z = sigma * epsilon + mu.
    For more stable training when integrating this function into a neural network, it helps to
    pass this function the log of the variance of the distribution from which to sample, rather
    than specifying the standard deviation directly.

    Inputs:
    - mu: Tensor of shape (N, Z) giving means
    - logvar: Tensor of shape (N, Z) giving log-variances

    Returns:
    - z: Estimated latent vectors, where z[i, j] is a random value sampled from a Gaussian with
        mean mu[i, j] and log-variance logvar[i, j].
    """
    z = None

    #####
    # TODO: Reparametrize by initializing epsilon as a normal distribution and scaling by
    #
    # posterior mu and sigma to estimate z
    #

    #####
    # Replace "pass" statement with your code
    epsilon = torch.randn(logvar.shape).to(device)

```

```

sigma = torch.exp(0.5 * logvar).to(device)
z = sigma * epsilon + mu

#####
#                                     END OF YOUR CODE
#

#####

return z

def loss_function(x_hat, x, mu, logvar):
    """
    Computes the negative variational lower bound loss term of the VAE (refer to formulation in
    notebook).
    Inputs:
    - x_hat: Reconstructed input data of shape (N, 1, H, W)
    - x: Input data for this timestep of shape (N, 1, H, W)
    - mu: Matrix representing estimated posterior mu (N, Z), with Z latent space dimension
    - logvar: Matrix representing estimated variance in log-space (N, Z), with Z latent space
    dimension
    Returns:
    - loss: Tensor containing the scalar loss for the negative variational lowerbound
    """
    loss = None

    #####
    # TODO: Compute negative variational lowerbound loss as described in the notebook
    #

    #####
    # Replace "pass" statement with your code
    reconstruction_loss = (
        nn.functional.binary_cross_entropy(x_hat, x, reduction="sum") / x_hat.shape[0]
    )

    kl_loss = -0.5 * torch.mean(
        torch.sum(1 + logvar - torch.square(mu) - torch.exp(logvar), dim=1)
    )

    loss = reconstruction_loss + kl_loss

    #####
    #                                     END OF YOUR CODE
    #

    #####

    return loss

```