

projects/project_3/code/model.py

```
## Building and training a bigram language model
```

```
import math
```

```
import torch
```

```
import torch.nn as nn
```

```
from config import BigramConfig
```

```
device = torch.device(
    "cuda"
    if torch.cuda.is_available()
    else "mps"
    if torch.mps.is_available()
    else "cpu"
)
```

```
class BigramLanguageModel(nn.Module):
```

```
    """
```

```
    Class definition for a simple bigram language model.
```

```
    """
```

```
def __init__(self, config: BigramConfig):
```

```
    """
```

```
    Initialize the bigram language model with the given configuration.
```

```
    Args:
```

```
    config : BigramConfig (Defined in config.py)
```

```
        Configuration object containing the model parameters.
```

```
    The model should have the following layers:
```

```
    1. An embedding layer that maps tokens to embeddings. (self.embeddings)
```

```
        You can use the Embedding layer from PyTorch.
```

```
    2. A linear layer that maps embeddings to logits. (self.linear) **set bias to
```

```
True**
```

```
    3. A dropout layer. (self.dropout)
```

```
    NOTE : PLEASE KEEP OF EACH LAYER AS PROVIDED BELOW TO FACILITATE TESTING.
```

```
    """
```

```
    super().__init__()
```

```
    # ===== TODO : START ===== #
```

```
    self.config = config
```

```
    self.embeddings = nn.Embedding(self.config.vocab_size, self.config.embed_dim)
```

```
    self.linear = nn.Linear(
        self.config.context_length * self.config.embed_dim,
        self.config.vocab_size,
        bias=True,
    )
```

```
    self.dropout = nn.Dropout(p=self.config.dropout)
```

```

# ===== TODO : END ===== #

self.apply(self._init_weights)

def forward(self, x):
    """
    Forward pass of the bigram language model.
    Args:
    x : torch.Tensor
        A tensor of shape (batch_size, 1) containing the input tokens.
    Output:
    torch.Tensor
        A tensor of shape (batch_size, vocab_size) containing the logits.
    """

# ===== TODO : START ===== #

embed = self.embeddings(x).squeeze(1) # (batch_size, embed_dim)
out = self.linear(embed) # (batch_size, vocab_size)
out = self.dropout(out) # (batch_size, vocab_size)

return out

# ===== TODO : END ===== #

def _init_weights(self, module):
    """
    Weight initialization for better convergence.
    NOTE : You do not need to modify this function.
    """

    if isinstance(module, nn.Linear):
        torch.nn.init.normal_(module.weight, mean=0.0, std=0.02)
        if module.bias is not None:
            torch.nn.init.zeros_(module.bias)
    elif isinstance(module, nn.Embedding):
        torch.nn.init.normal_(module.weight, mean=0.0, std=0.02)

def generate(self, context, max_new_tokens=100):
    """
    Use the model to generate new tokens given a context.
    We will perform multinomial sampling which is very similar to greedy sampling,
    but instead of taking the token with the highest probability, we sample the next
    token from a multinomial distribution.
    Remember in Bigram Language Model, we are only using the last token to predict
    the next token.
    You should sample the next token  $x_t$  from the distribution  $p(x_t | x_{\{t-1\}})$ .
    Args:
    context : List[int]

```

```

    A list of integers (tokens) representing the context.
max_new_tokens : int
    The maximum number of new tokens to generate.
Output:
List[int]
    A list of integers (tokens) representing the generated tokens.
"""

### ===== TODO : START ===== ###

f = torch.softmax
context = torch.tensor(context, device=device)
current_token = context[-1]
for _ in range(max_new_tokens):
    logits = self.forward(
        torch.tensor([current_token], device=device)
    ).squeeze()
    probabilities = f(logits, dim=0)
    current_token = torch.multinomial(probabilities, 1).to(device)
    context = torch.cat((context, current_token), dim=0)

return context
### ===== TODO : END ===== ###

```

```

class SingleHeadAttention(nn.Module):

```

```

    """

```

```

    Class definition for Single Head Causal Self Attention Layer.
    As in Attention is All You Need (https://arxiv.org/pdf/1706.03762)
    """

```

```

    """

```

```

    def __init__(

```

```

        self,
        input_dim,
        output_key_query_dim=None,
        output_value_dim=None,
        dropout=0.1,
        max_len=512,

```

```

    ):

```

```

        """

```

```

        Initialize the Single Head Attention Layer.

```

```

        The model should have the following layers:

```

1. A linear layer for key. (self.key) **set bias to False**
2. A linear layer for query. (self.query) **set bias to False**
3. A linear layer for value. (self.value) # **set bias to False**
4. A dropout layer. (self.dropout)
5. A causal mask. (self.causal_mask) This should be registered as a buffer.
 - You can use the torch.tril function to create a lower triangular matrix.
 - In the skeleton we use register_buffer to register the causal mask as a

```

        buffer.

```

This is typically used to register a buffer that should not to be considered a model parameter.

NOTE : Please make sure that the causal mask is upper triangular and not lower triangular (this helps in setting up the test cases,)

NOTE : PLEASE KEEP OF EACH LAYER AS PROVIDED BELOW TO FACILITATE TESTING.

```

"""
super().__init__()

self.input_dim = input_dim
if output_key_query_dim:
    self.output_key_query_dim = output_key_query_dim
else:
    self.output_key_query_dim = input_dim

if output_value_dim:
    self.output_value_dim = output_value_dim
else:
    self.output_value_dim = input_dim

causal_mask = None # You have to implement this, currently just a placeholder

# ===== TODO : START ===== #

self.key = nn.Linear(
    self.input_dim, self.output_key_query_dim, bias=False, device=device
)
self.query = nn.Linear(
    self.input_dim, self.output_key_query_dim, bias=False, device=device
)
self.value = nn.Linear(
    self.input_dim, self.output_value_dim, bias=False, device=device
)

self.dropout = nn.Dropout(p=dropout)

causal_mask = torch.ones((max_len, max_len), device=device)
causal_mask *= -float("inf")
causal_mask = torch.triu(causal_mask, 1)

# ===== TODO : END ===== #

self.register_buffer(
    "causal_mask", causal_mask
) # Registering as buffer to avoid backpropagation

def forward(self, x):
    """
    Forward pass of the Single Head Attention Layer.
    Args:
    x : torch.Tensor

```

A tensor of shape (batch_size, num_tokens, token_dim) containing the input tokens.

Output:

torch.Tensor

A tensor of shape (batch_size, num_tokens, output_value_dim) containing the output tokens.

Hint:

- You need to 'trim' the causal mask to the size of the input tensor.

"""

===== TODO : START =====

_, num_tokens, _ = x.shape

dk = self.output_key_query_dim

K = self.key(x) # (B, T, Dk)

Q = self.query(x) # (B, T, Dk)

V = self.value(x) # (B, T, Dv)

attention_scores = Q @ K.transpose(-2, -1) # (B, T, T)

attention_scores /= math.sqrt(dk) # (B, T, T)

mask = self.causal_mask[:num_tokens, :num_tokens] # (T, T)

mask = mask.unsqueeze(0) # not sure this is necessary # (1, T, T)

mask = mask.type(torch.bool)

mask = torch.where(mask, -torch.inf, 0)

attention_scores += mask # (B, T, T)

attention_weights = torch.softmax(attention_scores, dim=-1) # (B, T, T)

attention_weights = self.dropout(attention_weights) # (B, T, T)

out = attention_weights @ V # (B, T, Dv)

return out

===== TODO : END =====

class MultiHeadAttention(nn.Module):

"""

Class definition for Multi Head Attention Layer.

As in Attention is All You Need (<https://arxiv.org/pdf/1706.03762>)

"""

def __init__(self, input_dim, num_heads, dropout=0.1) → None:

"""

Initialize the Multi Head Attention Layer.

The model should have the following layers:

1. Multiple SingleHeadAttention layers. (self.head_{i}) Use setattr to dynamically set the layers.

2. A linear layer for output. (self.out) **set bias to True**

3. A dropout layer. (self.dropout) Apply dropout to the output of the out layer.

NOTE : PLEASE KEEP OF EACH LAYER AS PROVIDED BELOW TO FACILITATE TESTING.

"""

super().__init__()

self.input_dim = input_dim

self.num_heads = num_heads

self.head_dim = input_dim // num_heads

self.heads = []

for i in range(num_heads):

 head = SingleHeadAttention(
 input_dim=input_dim,
 output_key_query_dim=self.head_dim,
 output_value_dim=self.head_dim,
)

 setattr(self, f"head_{i}", head)

 self.heads.append(head)

self.out = nn.Linear(self.input_dim, self.input_dim)

self.dropout = nn.Dropout(p=dropout)

===== TODO : END =====

def forward(self, x):

"""

Forward pass of the Multi Head Attention Layer.

Args:

x : torch.Tensor

 A tensor of shape (batch_size, num_tokens, token_dim) containing the input tokens.

Output:

torch.Tensor

 A tensor of shape (batch_size, num_tokens, token_dim) containing the output tokens.

"""

===== TODO : START =====

head_outputs = []

for head in self.heads:

 head_outputs.append(head(x))

out = self.out(torch.cat(head_outputs, dim=-1))

out = self.dropout(out)

return out

===== TODO : END =====

```

class FeedForwardLayer(nn.Module):
    """
    Class definition for Feed Forward Layer.
    """

    def __init__(self, input_dim, feedforward_dim=None, dropout=0.1):
        """
        Initialize the Feed Forward Layer.
        The model should have the following layers:
        1. A linear layer for the feedforward network. (self.fc1) **set bias to True**
        2. A GELU activation function. (self.activation)
        3. A linear layer for the feedforward network. (self.fc2) ** set bias to True**
        4. A dropout layer. (self.dropout)
        NOTE : PLEASE KEEP OF EACH LAYER AS PROVIDED BELOW TO FACILITATE TESTING.
        """
        super().__init__()

        if feedforward_dim is None:
            feedforward_dim = input_dim * 4

        # ===== TODO : START ===== #

        self.fc1 = nn.Linear(input_dim, feedforward_dim, bias=True)
        self.activation = nn.GELU()
        self.fc2 = nn.Linear(feedforward_dim, input_dim, bias=True)
        self.dropout = nn.Dropout(dropout)

        # ===== TODO : END ===== #

    def forward(self, x):
        """
        Forward pass of the Feed Forward Layer.
        Args:
        x : torch.Tensor
            A tensor of shape (batch_size, num_tokens, token_dim) containing the input
tokens.
        Output:
        torch.Tensor
            A tensor of shape (batch_size, num_tokens, token_dim) containing the output
tokens.
        """

        ### ===== TODO : START ===== ###

        out = self.fc1(x)
        out = self.activation(out)
        out = self.fc2(out)
        out = self.dropout(out)
        return out

```

```
### ===== TODO : END ===== ###
```

```
class LayerNorm(nn.Module):
    """
    LayerNorm module as in the paper https://arxiv.org/abs/1607.06450
    Note : Variance computation is done with biased variance.
    Hint :
    - You can use torch.var and specify whether to use biased variance or not.
    """

    def __init__(self, normalized_shape, eps=1e-05, elementwise_affine=True) → None:
        super().__init__()

        self.normalized_shape = (normalized_shape,)
        self.eps = eps
        self.elementwise_affine = elementwise_affine

        if elementwise_affine:
            self.gamma = nn.Parameter(torch.ones(tuple(self.normalized_shape)))
            self.beta = nn.Parameter(torch.zeros(tuple(self.normalized_shape)))

    def forward(self, input):
        """
        Forward pass of the LayerNorm Layer.
        Args:
        input : torch.Tensor
            A tensor of shape (batch_size, num_tokens, token_dim) containing the input
tokens.
        Output:
        torch.Tensor
            A tensor of shape (batch_size, num_tokens, token_dim) containing the output
tokens.
        """

        mean = None
        var = None
        # ===== TODO : START ===== #

        mean = torch.mean(input, dim=2, keepdim=True)
        var = torch.var(input, dim=2, keepdim=True, correction=0)

        # ===== TODO : END ===== #

        if self.elementwise_affine:
            return (
                self.gamma * (input - mean) / torch.sqrt((var + self.eps)) + self.beta
            )
        else:
            return (input - mean) / torch.sqrt((var + self.eps))
```



```

class TransformerLayer(nn.Module):
    """
    Class definition for a single transformer layer.
    """

    def __init__(self, input_dim, num_heads, feedforward_dim=None):
        super().__init__()
        """
        Initialize the Transformer Layer.
        We will use prenorm layer where we normalize the input before applying the
        attention and feedforward layers.
        The model should have the following layers:
        1. A LayerNorm layer. (self.norm1)
        2. A MultiHeadAttention layer. (self.attention)
        3. A LayerNorm layer. (self.norm2)
        4. A FeedForwardLayer layer. (self.feedforward)
        NOTE : PLEASE KEEP OF EACH LAYER AS PROVIDED BELOW TO FACILITATE TESTING.
        """

        # ===== TODO : START ===== #

        self.norm1 = LayerNorm(normalized_shape=input_dim)
        self.attention = MultiHeadAttention(input_dim=input_dim, num_heads=num_heads)
        self.norm2 = LayerNorm(normalized_shape=input_dim)
        self.feedforward = FeedForwardLayer(
            input_dim=input_dim, feedforward_dim=feedforward_dim
        )

        # ===== TODO : END ===== #

    def forward(self, x):
        """
        Forward pass of the Transformer Layer.
        Args:
        x : torch.Tensor
            A tensor of shape (batch_size, num_tokens, token_dim) containing the input
tokens.
        Output:
        torch.Tensor
            A tensor of shape (batch_size, num_tokens, token_dim) containing the output
tokens.
        """

        # ===== TODO : START ===== #

        intermediate = self.norm1(x)
        intermediate = self.attention(intermediate)
        intermediate += x

```

```

    out = self.norm2(intermediate)
    out = self.feedforward(out)
    out += intermediate

```

```

    return out

```

```

# ===== TODO : END ===== #

```

```

class MiniGPT(nn.Module):

```

```

    """

```

```

    Putting it all together: GPT model

```

```

    """

```

```

    def __init__(self, config) → None:

```

```

        super().__init__()

```

```

        """

```

```

        Putting it all together: our own GPT model!

```

```

        Initialize the MiniGPT model.

```

```

        The model should have the following layers:

```

1. An embedding layer that maps tokens to embeddings. (self.vocab_embedding)
2. A positional embedding layer. (self.positional_embedding) We will use learnt positional embeddings.
3. A dropout layer for embeddings. (self.embed_dropout)
4. Multiple TransformerLayer layers. (self.transformer_layers)
5. A LayerNorm layer before the final layer. (self.prehead_norm)
6. Final language Modelling head layer. (self.head) We will use weight tying (<https://paperswithcode.com/method/weight-tying>) and set the weights of the head layer to be the same as the vocab_embedding layer.

```

        NOTE: You do not need to modify anything here.

```

```

        """

```

```

        self.context_length = config.context_length

```

```

        self.vocab_embedding = nn.Embedding(config.vocab_size, config.embed_dim)

```

```

        self.positional_embedding = nn.Embedding(
            config.context_length, config.embed_dim
        )

```

```

        self.embed_dropout = nn.Dropout(config.embed_dropout)

```

```

        self.transformer_layers = nn.ModuleList(

```

```

            [

```

```

                TransformerLayer(

```

```

                    config.embed_dim, config.num_heads, config.feedforward_size

```

```

                )

```

```

                for _ in range(config.num_layers)

```

```

            ]

```

```

        )

```

```

        # prehead layer norm

```

```

        self.prehead_norm = LayerNorm(config.embed_dim)

```

```

self.head = nn.Linear(
    config.embed_dim, config.vocab_size
) # Language modelling head

if config.weight_tie:
    self.head.weight = self.vocab_embedding.weight

# precreate positional indices for the positional embedding
pos = torch.arange(0, config.context_length, dtype=torch.long)
self.register_buffer("pos", pos, persistent=False)

self.apply(self._init_weights)

def forward(self, x):
    """
    Forward pass of the MiniGPT model.
    Remember to add the positional embeddings to your input token!!
    Args:
    x : torch.Tensor
        A tensor of shape (batch_size, seq_len) containing the input tokens.
    Output:
    torch.Tensor
        A tensor of shape (batch_size, seq_len, vocab_size) containing the logits.
    Hint:
    - You may need to 'trim' the positional embedding to match the input sequence
length
    """

    ### ===== TODO : START ===== ###

    B, T = x.shape

    # embeddings
    token_embeds = self.vocab_embedding(x)
    position_embeds = self.positional_embedding(self.pos[:T])
    position_embeds = position_embeds.unsqueeze(dim=0)
    x = token_embeds + position_embeds
    x = self.embed_dropout(x)

    # attention layers
    for layer in self.transformer_layers:
        x = layer(x)

    # language modeling head
    x = self.prehead_norm(x)
    logits = self.head(x)

    return logits

```

```

### ===== TODO : END ===== ###

def _init_weights(self, module):
    """
    Weight initialization for better convergence.
    NOTE : You do not need to modify this function.
    """

    if isinstance(module, nn.Linear):
        if module._get_name() == "fc2":
            # GPT-2 style FFN init
            torch.nn.init.normal_(
                module.weight, mean=0.0, std=0.02 / math.sqrt(2 * self.num_layers)
            )
        else:
            torch.nn.init.normal_(module.weight, mean=0.0, std=0.02)
        if module.bias is not None:
            torch.nn.init.zeros_(module.bias)
    elif isinstance(module, nn.Embedding):
        torch.nn.init.normal_(module.weight, mean=0.0, std=0.02)

def generate(self, context, max_new_tokens=100):
    """
    Use the model to generate new tokens given a context.
    Hint:
    - This should be similar to the Bigram Language Model, but you will use the
    entire context to predict the next token.
    Instead of sampling from the distribution  $p(x_t | x_{t-1})$ ,
    you will sample from the distribution  $p(x_t | x_{t-1}, x_{t-2}, \dots, x_{t-n})$ 
    ,
    where n is the context length.
    - When decoding for the next token, you should use the logits of the last token
    in the input sequence.
    """

    ### ===== TODO : START ===== ###

    f = torch.softmax
    context = torch.tensor(context, device=device).unsqueeze(0)
    for i in range(max_new_tokens):
        current_context = context[:, -(self.context_length) :]
        logits = self(current_context)
        next_token_logits = logits[:, -1, :]
        next_token_probabilities = f(next_token_logits, dim=-1)
        current_token = torch.multinomial(
            input=next_token_probabilities, num_samples=1
        ).to(device)
        context = torch.cat((context, current_token), dim=1)

    return context

```

```
### ===== TODO : END ===== ###
```

```
class SingleHeadGeneralAttention(nn.Module):
    def __init__(
        self,
        input_dim,
        output_key_query_dim=None,
        output_value_dim=None,
        dropout=0.1,
        max_len=512,
    ):
        super().__init__()

        self.input_dim = input_dim
        if output_key_query_dim:
            self.output_key_query_dim = output_key_query_dim
        else:
            self.output_key_query_dim = input_dim

        if output_value_dim:
            self.output_value_dim = output_value_dim
        else:
            self.output_value_dim = input_dim

        causal_mask = None # You have to implement this, currently just a placeholder

        self.key = nn.Linear(
            self.input_dim, self.output_key_query_dim, bias=False, device=device
        )
        self.query = nn.Linear(
            self.input_dim, self.output_key_query_dim, bias=False, device=device
        )
        self.value = nn.Linear(
            self.input_dim, self.output_value_dim, bias=False, device=device
        )

        self.dropout = nn.Dropout(p=dropout)

        causal_mask = torch.ones((max_len, max_len), device=device)
        causal_mask *= -float("inf")
        causal_mask = torch.triu(causal_mask, 1)

        self.register_buffer(
            "causal_mask", causal_mask
        ) # Registering as buffer to avoid backpropagation

    def forward(self, x_query, x_key, x_value):
        _, num_tokens, _ = x_query.shape
```

```

dk = self.output_key_query_dim
K = self.key(x_key) # (B, T, Dk)
Q = self.query(x_query) # (B, T, Dk)
V = self.value(x_value) # (B, T, Dv)

attention_scores = Q @ K.transpose(-2, -1) # (B, T, T)
attention_scores /= math.sqrt(dk) # (B, T, T)

mask = self.causal_mask[:num_tokens, :num_tokens] # (T, T)
mask = mask.unsqueeze(0)
mask = mask.type(torch.bool)
mask = torch.where(mask, -torch.inf, 0)

attention_scores += mask # (B, T, T)

attention_weights = torch.softmax(attention_scores, dim=-1) # (B, T, T)
attention_weights = self.dropout(attention_weights) # (B, T, T)

out = attention_weights @ V # (B, T, Dv)
return out

```

```

class MultiHeadGeneralAttention(nn.Module):
    def __init__(self, input_dim, num_heads, dropout=0.1) → None:
        super().__init__()

        self.input_dim = input_dim
        self.num_heads = num_heads
        self.head_dim = input_dim // num_heads

        self.heads = []
        for i in range(num_heads):
            head = SingleHeadGeneralAttention(
                input_dim=input_dim,
                output_key_query_dim=self.head_dim,
                output_value_dim=self.head_dim,
            )
            setattr(self, f"head_{i}", head)
            self.heads.append(head)

        self.out = nn.Linear(self.input_dim, self.input_dim)
        self.dropout = nn.Dropout(p=dropout)

    def forward(self, x_query, x_key, x_value):
        head_outputs = []
        for head in self.heads:
            head_outputs.append(head(x_query, x_key, x_value))

        out = self.out(torch.cat(head_outputs, dim=-1))
        out = self.dropout(out)

```

```
return out
```

```
class GeneralTransformerLayer(nn.Module):
    def __init__(self, input_dim, num_heads, feedforward_dim=None):
        super().__init__()

        self.norm1 = LayerNorm(normalized_shape=input_dim)
        self.attention = MultiHeadGeneralAttention(
            input_dim=input_dim, num_heads=num_heads
        )
        self.norm2 = LayerNorm(normalized_shape=input_dim)
        self.feedforward = FeedForwardLayer(
            input_dim=input_dim, feedforward_dim=feedforward_dim
        )

    def forward(self, x):
        intermediate = self.norm1(x)
        intermediate = self.attention(intermediate)
        intermediate += x
        out = self.norm2(intermediate)
        out = self.feedforward(out)
        out += intermediate

        return out
```

```
class Encoder(nn.Module):
    def __init__(self, config) → None:
        super().__init__()

        self.context_length = config.context_length
        self.vocab_embedding = nn.Embedding(config.vocab_size, config.embed_dim)
        self.positional_embedding = nn.Embedding(
            config.context_length, config.embed_dim
        )
        self.embed_dropout = nn.Dropout(config.embed_dropout)

        self.transformer_layers = nn.ModuleList(
            [
                TransformerLayer(
                    config.embed_dim, config.num_heads, config.feedforward_size
                )
                for _ in range(config.num_layers)
            ]
        )

        # precreate positional indices for the positional embedding
        pos = torch.arange(0, config.context_length, dtype=torch.long)
        self.register_buffer("pos", pos, persistent=False)
```

```

        self.apply(self._init_weights)

    def forward(self, x):
        B, T = x.shape

        # embeddings
        token_embeds = self.vocab_embedding(x)
        position_embeds = self.positional_embedding(self.pos[:T])
        position_embeds = position_embeds.unsqueeze(dim=0)
        x = token_embeds + position_embeds
        x = self.embed_dropout(x)

        # attention layers
        for layer in self.transformer_layers:
            x = layer(x)

        return x

class Decoder(nn.Module):
    def __init__(self, config) → None:
        super().__init__()

        self.context_length = config.context_length
        self.vocab_embedding = nn.Embedding(config.vocab_size, config.embed_dim)
        self.positional_embedding = nn.Embedding(
            config.context_length, config.embed_dim
        )
        self.embed_dropout = nn.Dropout(config.embed_dropout)

        self.decoder_layers = nn.ModuleList(
            [
                TransformerLayer(
                    config.embed_dim, config.num_heads, config.feedforward_size
                )
                for _ in range(config.num_layers)
            ]
        )

        # prehead layer norm
        self.prehead_norm = LayerNorm(config.embed_dim)

        self.head = nn.Linear(
            config.embed_dim, config.vocab_size
        ) # Language modelling head

        if config.weight_tie:
            self.head.weight = self.vocab_embedding.weight

```



```

# precreate positional indices for the positional embedding
pos = torch.arange(0, config.context_length, dtype=torch.long)
self.register_buffer("pos", pos, persistent=False)

self.apply(self._init_weights)

def forward(self, x, encoder_output):
    B, T = x.shape

    # embeddings
    token_embeds = self.vocab_embedding(x)
    position_embeds = self.positional_embedding(self.pos[:T])
    position_embeds = position_embeds.unsqueeze(dim=0)
    x = token_embeds + position_embeds
    x = self.embed_dropout(x)

    # attention layers
    for layer in self.decoder_layers:
        x = layer(x_query=x, x_key=encoder_output, x_value=encoder_output)

    # language modeling head
    x = self.prehead_norm(x)
    logits = self.head(x)

    return logits

def _init_weights(self, module):
    if isinstance(module, nn.Linear):
        if module._get_name() == "fc2":
            # GPT-2 style FFN init
            torch.nn.init.normal_(
                module.weight, mean=0.0, std=0.02 / math.sqrt(2 * self.num_layers)
            )
        else:
            torch.nn.init.normal_(module.weight, mean=0.0, std=0.02)
        if module.bias is not None:
            torch.nn.init.zeros_(module.bias)
    elif isinstance(module, nn.Embedding):
        torch.nn.init.normal_(module.weight, mean=0.0, std=0.02)

def generate(self, context, max_new_tokens=100):
    f = torch.softmax
    context = torch.tensor(context, device=device).unsqueeze(0)
    for i in range(max_new_tokens):
        current_context = context[:, -(self.context_length) :]
        logits = self(current_context)
        next_token_logits = logits[:, -1, :]
        next_token_probabilities = f(next_token_logits, dim=-1)
        current_token = torch.multinomial(
            input=next_token_probabilities, num_samples=1

```

```
    ).to(device)  
    context = torch.cat((context, current_token), dim=1)  
  
    return context
```