

projects/project_3/code/train.py

```
"""
```

```
Training file for the models we implemented
```

```
"""
```

```
from pathlib import Path
```

```
import torch
```

```
import torch.nn.utils
```

```
import torch.nn as nn
```

```
from torch.utils.data import DataLoader
```

```
import wandb
```

```
from tqdm import tqdm
```

```
from model import BigramLanguageModel, MiniGPT
```

```
from dataset import TinyStoriesDataset
```

```
from config import BigramConfig, MiniGPTConfig
```

```
def solver(model_name):
```

```
    # Initialize the model
```

```
    if model_name == "bigram":
```

```
        config = BigramConfig(to_log=True, log_interval=10_000)
```

```
        model = BigramLanguageModel(config)
```

```
    elif model_name == "minigpt":
```

```
        config = MiniGPTConfig(
```

```
            to_log=True, save_iterations=100_000, log_interval=100_000
```

```
        )
```

```
        model = MiniGPT(config)
```

```
    else:
```

```
        raise ValueError("Invalid model name")
```

```
    # Load the dataset
```

```
    train_dataset = TinyStoriesDataset(
```

```
        config.path_to_data,
```

```
        mode="train",
```

```
        context_length=config.context_length,
```

```
    )
```

```
    eval_dataset = TinyStoriesDataset(
```

```
        config.path_to_data, mode="test", context_length=config.context_length
```

```
    )
```

```
    # Create the dataloaders
```

```
    train_dataloader = DataLoader(
```

```
        train_dataset, batch_size=config.batch_size, pin_memory=True
```

```
    )
```

```
    eval_dataloader = DataLoader(
```

```
        eval_dataset, batch_size=config.batch_size, pin_memory=True
```

```
    )
```

```

# Set the device
device = torch.device(
    "cuda"
    if torch.cuda.is_available()
    else "mps"
    if torch.mps.is_available()
    else "cpu"
)

# Print number of parameters in the model
def count_parameters(model):
    return sum(p.numel() for p in model.parameters() if p.requires_grad)

print("number of trainable parameters: %.2fM" % (count_parameters(model) / 1e6,))

# Initialize wandb if you want to use it
if config.to_log:
    wandb.login()
    wandb.init(
        project="dl2_proj3",
        config={
            "learning_rate": config.learning_rate,
            "architecture": model_name,
            "dataset": "TinyStories",
        },
    )

# Create the save path if it does not exist
if not Path.exists(config.save_path):
    Path.mkdir(config.save_path, parents=True, exist_ok=True)

### ===== START OF YOUR CODE ===== ###
"""
You are required to implement the training loop for the model.
The code below is a skeleton for the training loop, for your reference.
You can fill in the missing parts or completely set it up from scratch.
Please keep the following in mind:
- You will need to define an appropriate loss function for the model.
- You will need to define an optimizer for the model.
- You are required to log the loss (either on wandb or any other logger you prefer)
every `config.log_interval` iterations.
- It is recommended that you save the model weights every `config.save_iterations`
iterations. You can also just save the model with the best training loss.
NOTE :
- Please check the config file to see the different configurations you can set for
the model.
- The MiniGPT config has params that you do not need to use, these were added to
scale the model but are
not a required part of the assignment.

```

- Feel free to experiment with the parameters and I would be happy to talk to you about them if interested.

"""

===== TODO : START =====

Define the loss function

loss = nn.CrossEntropyLoss()

Define the optimizer

optimizer = torch.optim.Adam(params=model.parameters(), lr=config.learning_rate)

===== TODO : END =====

if config.scheduler:

 scheduler = torch.optim.lr_scheduler.CosineAnnealingWarmRestarts(
 optimizer, T_0=2000, T_mult=2
)

model.train()

model.to(device)

max_iter_eval = len(eval_data_loader) // 500

for i, (context, target) in tqdm(enumerate(train_data_loader)):

 context = context.to(device)

 target = target.to(device)

 train_loss = 0 # You can use this variable to store the training loss for the current iteration

===== TODO : START =====

Do the forward pass, compute the loss, do the backward pass, and update the weights with the optimizer.

model.zero_grad()

logits = model(context)

target = target.long()

if model_name == "bigram":

 target = target.squeeze()

elif model_name == "minigpt":

 B, T, _ = logits.shape

 logits = logits.reshape(B * T, -1)

 target = target.reshape(B * T)

batch_loss = loss(logits, target)

batch_loss.backward()

optimizer.step()

train_loss += batch_loss.item()

```

### ===== TODO : END ===== ###

if config.scheduler:
    scheduler.step()

del context, target # Clear memory

if i % config.log_interval == 0:
    model.eval()
    eval_loss = 0 # You can use this variable to store the evaluation loss for
the current iteration

### ===== TODO : START ===== ###
# Compute the evaluation loss on the eval dataset.

with torch.no_grad():
    for j, (context, target) in enumerate(eval_data_loader):
        context = context.to(device)
        target = target.to(device)

        target = target.long()
        logits = model(context)

        if model_name == "bigram":
            target = target.squeeze()

        elif model_name == "minigpt":
            logits = logits.reshape(B * T, -1)
            target = target.reshape(B * T)

        batch_loss = loss(logits, target)
        eval_loss += batch_loss.item()
        del context, target # Clear memory

        if j > max_iter_eval:
            break

### ===== TODO : END ===== ###

eval_loss /= max_iter_eval
print(
    f"Iteration {i}\nTrain Loss: {train_loss}\nAverage eval loss: {eval_loss}
",
)

if config.to_log:
    wandb.log(
        {
            "Train Loss": train_loss,
            "Eval Loss": eval_loss,

```

```
        }  
    )  
  
    model.train()  
  
    # Save the model every config.save_iterations  
    if i % config.save_iterations == 0:  
        torch.save(  
            {  
                "model_state_dict": model.state_dict(),  
                "optimizer_state_dict": optimizer.state_dict(),  
                "train_loss": train_loss,  
                "eval_loss": eval_loss,  
                "iteration": i,  
            },  
            config.save_path / f"mini_model_checkpoint_{i}.pt",  
        )  
  
    if i > config.max_iter:  
        break
```